

RTX 5090 sm_120 CUDA Compatibility Guide for Docker ML Workloads

Executive Summary

The NVIDIA RTX 5090 with Blackwell architecture (compute capability 12.0/sm_120) launched January 30, 2025, [\(NVIDIA +3\)](#) presenting significant compatibility challenges for ML workloads in Docker containers. [\(MIT HAN Lab +4\)](#) **Critical requirement: CUDA 12.8+ and specialized configurations are mandatory**, as current stable ML frameworks lack native sm_120 support. [\(Super User\)](#) This guide provides comprehensive solutions for running embedding models like BGE-M3 and Multilingual E5 Large effectively on the RTX 5090's 32GB GDDR7 memory. [\(Markaicode +4\)](#)

1. CUDA Toolkit Requirements for sm_120 Architecture

Minimum and Recommended Versions

- **Minimum:** CUDA Toolkit 12.8 GA (first version with native sm_120 support) [\(GitHub +4\)](#)
- **Recommended Production:** CUDA 12.9 Update 1 or CUDA 13.0 Update 1
- **Driver Requirements:**
 - Minimum: 570.xxx.xx [\(GitHub +2\)](#)
 - Current stable: 572.24 (includes critical fixes) [\(Graphics Card Hub\)](#)
 - Linux: ≥575.57.08 for CUDA 12.9 [\(nvidia\)](#)
 - Windows: ≥576.57 for CUDA 12.9 [\(nvidia\)](#)

Compilation Flags for sm_120

```
bash

-gencode=arch=compute_80,code=sm_80 \
-gencode=arch=compute_89,code=sm_89 \
-gencode=arch=compute_90,code=sm_90 \
-gencode=arch=compute_120,code=sm_120 \
-gencode=arch=compute_120,code=compute_120
```

Known Issues (January 2025)

- **PCIe Gen 5 compatibility:** Set motherboard to PCIe Gen 4 mode if experiencing issues (minimal 1% performance impact) [\(Tom's Hardware\)](#) [\(Graphics Card Hub\)](#)

- **Driver installation black screens:** Resolved with firmware updates in 572.24 [PC Gamer](#)
[Graphics Card Hub](#)
- **GPU virtualization freezes:** Workaround requires disabling modesetting [MIT HAN Lab](#) [VideoCardz](#)

2. Docker Configuration Requirements

Essential Software Stack

- **Docker:** Version 19.03+ (for --gpu flag support) [Stack Overflow](#) [NVIDIA](#)
- **NVIDIA Container Toolkit:** Version 1.17.8+ [NVIDIA](#) (critical for sm_120)
- **Base OS:** Ubuntu 20.04+, Ubuntu 24.04 recommended

Installation and Configuration

```
bash

# Install NVIDIA Container Toolkit
curl -fsSL https://nvidia.github.io/libnvidia-container/gpgkey | \
  sudo gpg --dearmor -o /usr/share/keyrings/nvidia-container-toolkit-keyring.gpg

curl -s -L https://nvidia.github.io/libnvidia-container/stable/deb/nvidia-container-toolkit.list | \
  sed 's#deb https://#deb [signed-by=/usr/share/keyrings/nvidia-container-toolkit-keyring.gpg] https://#g' | \
  sudo tee /etc/apt/sources.list.d/nvidia-container-toolkit.list

sudo apt-get update
sudo apt-get install -y nvidia-container-toolkit=1.17.8-1

# Configure Docker runtime
sudo nvidia-ctk runtime configure --runtime=docker
sudo systemctl restart docker
```

Essential Docker Run Parameters

```
bash
```

```
docker run -it --rm \
  --gpus all \
  --ipc=host \
  --ulimit memlock=-1 \
  --ulimit stack=67108864 \
  --shm-size=32g \
  -e PYTORCH_CUDA_ALLOC_CONF="max_split_size_mb:512,garbage_collection_threshold:0.8" \
  -e TORCH_CUDA_ARCH_LIST="12.0" \
  -v $(pwd):/workspace \
  [base-image]
```

3. PyTorch and TensorFlow Compatibility

PyTorch Status

Current State: PyTorch 2.7.0 provides Linux support via CUDA 12.8 wheels; [GitHub](#) Windows lacks stable support [PyTorch](#) [SaladCloud](#)

Installation Options:

```
bash

# Linux (Stable)
pip install torch==2.7.0 --index-url https://download.pytorch.org/whl/cu128

# Nightly (partial Windows support)
pip install --pre torch --index-url https://download.pytorch.org/whl/nightly/cu128

# Compatibility Layer (immediate workaround)
pip install rtx50-compat

# In Python:
import rtx50_compat # MUST import before PyTorch
import torch
```

Build from Source (if needed):

```
bash

export TORCH_CUDA_ARCH_LIST="12.0"
export FORCE_CUDA=1
export CUDA_HOME=/usr/local/cuda
export CUDACXX=$CUDA_HOME/bin/nvcc
python setup.py install
```

TensorFlow Status

Critical: No stable TensorFlow release supports sm_120. Nightly builds (2.20.0.dev) have significant issues including PTX compilation failures and runtime errors. [GitHub +3](#) Not recommended for production use.

4. Transformer Model Issues and Solutions

Common Errors and Fixes

Error: `"CUDA error: no kernel image is available for execution on the device"` [GitHub +7](#) **Solution:** Use PyTorch 2.7+ or rtx50-compat package [GitHub](#) [GitHub](#)

Flash Attention Compatibility

Community Build (Recommended):

- Source: [loscrossos/lib_flashattention v2.7.4.post1_crossos00](#) [GitHub](#) [github](#)
- Supports all RTX 50-series cards [GitHub](#)
- Compatible with PyTorch 2.7.0, Python 3.12 [github](#)

Alternative: PyTorch native SDPA with PR [#145602](#) adds sm_120 support [Hugging Face](#) [GitHub](#)

BGE-M3 and E5 Large Optimizations

```
python

# Enable Flash Attention 2
model.config.use_flash_attention_2 = True

# Optimal batch sizes for BGE-M3 (567M params)
# Sequence length → Batch size
# 512 → 64
# 1024 → 32
# 2048 → 16
# 4096 → 8

# Memory requirements
# BGE-M3: 2-3GB VRAM (inference), 8-12GB (fine-tuning)
# E5 Large: Use proper prefixes ("query:", "passage:")
```

Mixed Precision Recommendations

- Use BF16 over FP16 for numerical stability on Blackwell

- Enable TF32: `torch.backends.cuda.matmul.allow_tf32 = True` [Hugging Face](#)
- Leverage FP4 support for 2x performance over FP8 [MIT HAN Lab +2](#)

5. Optimal Docker Base Images

Recommended Pre-built Images

Best Option: `dconsorte/pytorch-tensorflow-gpu:latest`

- Ubuntu 24.04 + Python 3.11 [GitHub](#) [github](#)
- CUDA 12.8.1 + cuDNN 9.x [GitHub](#) [github](#)
- PyTorch 2.6.0.dev with Blackwell support [GitHub](#) [github](#)
- TensorFlow 2.19.0.dev (limited functionality) [GitHub](#) [github](#)
- Pre-configured Jupyter Lab [GitHub](#) [github](#)

Official CUDA Images (requires framework installation):

```
bash
```

```
nvidia/cuda:12.8.1-base-ubuntu24.04  
nvidia/cuda:12.8.1-runtime-ubuntu24.04  
nvidia/cuda:12.8.1-devel-ubuntu24.04
```

Custom Dockerfile for Production

```
dockerfile
```

```
FROM nvcr.io/nvidia/pytorch:25.04-py3
```

```
# Install CUDA 12.8+ for sm_120
```

```
RUN wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2404/x86_64/cuda-keyring_1.1-1_all.deb && \
    dpkg -i cuda-keyring_1.1-1_all.deb && \
    apt-get update && \
    apt-get install -y cuda-toolkit-12-8
```

```
# Set RTX 5090 environment variables
```

```
ENV TORCH_CUDA_ARCH_LIST="12.0"
```

```
ENV FORCE_CUDA=1
```

```
ENV TORCH_USE_CUDA_DSA=1
```

```
ENV CUDA_HOME=/usr/local/cuda
```

```
ENV CUDACXX=$CUDA_HOME/bin/nvcc
```

```
ENV PATH=$CUDA_HOME/bin:$PATH
```

```
# Install PyTorch with CUDA 12.8
```

```
RUN pip install torch==2.7.0 torchvision torchaudio --index-url https://download.pytorch.org/whl/cu128
```

```
# Install ML libraries
```

```
RUN pip install transformers accelerate datasets flash-attn sentence-transformers
```

```
# Enable TF32 optimization
```

```
RUN echo "import torch; torch.backends.cuda.matmul.allow_tf32 = True; torch.backends.cudnn.allow_tf32 = True"
```

```
WORKDIR /workspace
```

6. Memory Management for 32GB VRAM

Optimal Configuration

```
bash
```

```
export PYTORCH_CUDA_ALLOC_CONF="max_split_size_mb:512,roundup_power2_divisions:16,garbage_collection_t
```

Markaicode

Python Memory Settings

```
python
```

```

import torch

# Reserve 30GB for models (keep 2GB for system)
torch.cuda.set_per_process_memory_fraction(0.94)
torch.cuda.empty_cache()

# Enable gradient checkpointing for large models
model.gradient_checkpointing_enable()

# Dynamic batch size discovery
def find_optimal_batch_size(model, max_memory_gb=28):
    batch_sizes = [1, 2, 4, 8, 16, 32, 64, 128]
    for batch_size in batch_sizes:
        try:
            # Test forward pass
            dummy_input = torch.randn(batch_size, 512).cuda()
            _ = model(dummy_input)

            memory_used = torch.cuda.memory_allocated() / 1e9
            if memory_used > max_memory_gb:
                return batch_size // 2
            torch.cuda.empty_cache()
        except RuntimeError:
            return batch_size // 2
    return batch_sizes[-1]

```

Model Size Guidelines

Model Size	FP32	FP16	INT8	Recommended Batch Size (2048 tokens)
7B params	28GB	14GB	7GB	16
13B params	52GB	26GB	13GB	8
70B params	280GB	140GB	70GB	Requires quantization + offloading

Memory Profiling

```
python
```

```
# PyTorch profiler
with torch.profiler.profile(
    activities=[torch.profiler.ProfilerActivity.CUDA],
    profile_memory=True,
    record_shapes=True
) as prof:
    model(inputs)
prof.export_chrome_trace("memory_trace.json")

# Real-time monitoring
nvidia-smi -l 1 --query-gpu=memory.used,memory.free --format=csv
```

7. Latest NVIDIA Updates (January 2025)

Current Status

- **RTX 5090 Launch:** January 30, 2025 at \$1,999 [NVIDIA +3](#)
- **Architecture:** Blackwell (GB202), 21,760 CUDA cores, [Vast.ai](#) 575W TDP [Runpod +4](#)
- **Driver Updates:**
 - 572.16: Launch driver
 - 572.24: Critical stability fixes (current recommended)

Resolution Timeline

- **PyTorch Windows stable:** Expected in 2.7.1 or 2.8.0 (Q1-Q2 2025)
- **TensorFlow stable:** Timeline unclear, significant development needed
- **Flash Attention 3:** Native Blackwell support in development [GitHub](#)

Performance Expectations

- **LLM Inference:** 180-250 tokens/s (8B model), 120-180 tokens/s (13B model) [HostKey](#) [GitHub](#)
- **Training:** ~2,000 samples/hour (7B FP16), ~1,200 samples/hour (13B FP16)
- **vs RTX 4090:** 20-50% improvement in transformer workloads [HostKey +4](#)

Quick Start Commands

1. Verify RTX 5090 Setup

```
bash
```



```
# Check GPU detection
nvidia-smi
docker run --rm --gpus all nvidia/cuda:12.8.1-base-ubuntu24.04 nvidia-smi

# Test PyTorch compatibility
docker run --rm --gpus all -it dconsorte/pytorch-tensorflow-gpu:latest python -c \
    "import rtx50_compat; import torch; print(f'CUDA available: {torch.cuda.is_available()}'); \
    print(f'Device: {torch.cuda.get_device_name(0)}')"
```

2. Run BGE-M3 or E5 Large

```
python

# In container with proper setup
import rtx50_compat # If using compatibility layer
import torch
from sentence_transformers import SentenceTransformer

# Load model
model = SentenceTransformer('BAAI/bge-m3')
model = model.cuda()
model.config.use_flash_attention_2 = True

# Configure memory
torch.cuda.set_per_process_memory_fraction(0.94)
os.environ['PYTORCH_CUDA_ALLOC_CONF'] = 'max_split_size_mb:512,garbage_collection_threshold:0.8'

# Process embeddings
sentences = ["example text"] * 16 # Batch of 16
embeddings = model.encode(sentences, batch_size=16, convert_to_tensor=True)
```

3. Docker Compose Configuration

```
yaml
```

```
version: '3.8'
services:
  ml-training:
    image: dconsorte/pytorch-tensorflow-gpu:latest
    deploy:
      resources:
        reservations:
          devices:
            - driver: nvidia
              count: 1
              capabilities: [gpu]
    environment:
      - CUDA_VISIBLE_DEVICES=0
      - TORCH_CUDA_ARCH_LIST=12.0
      - PYTORCH_CUDA_ALLOC_CONF=max_split_size_mb:512,garbage_collection_threshold:0.8
    volumes:
      - ./workspace:/workspace
      - ./models:/models
    shm_size: '32gb'
    ulimits:
      memlock: -1
      stack: 67108864
```

Critical Takeaways

1. Always use CUDA 12.8+ and PyTorch 2.7+ for RTX 5090 support ([GitHub](#)) ([PyTorch](#))
2. Import rtx50-compat before PyTorch as immediate workaround ([PyPI](#)) ([GitHub](#))
3. Configure Docker with 32GB shared memory for large models
4. Use BF16 mixed precision for optimal performance/stability balance
5. Install community Flash Attention builds for transformer acceleration
6. Set PCIe to Gen 4 mode if experiencing compatibility issues
7. Monitor memory continuously and use gradient checkpointing for large models ([Hugging Face](#))
8. Expect 20-50% performance gains over RTX 4090 once properly configured ([HostKey](#))

[Computer Vision Lab](#)

This configuration enables full utilization of the RTX 5090's capabilities for ML workloads while navigating current compatibility limitations. ([Markaicode](#)) ([GitHub](#)) Regular driver and framework updates throughout 2025 will progressively improve native support.